

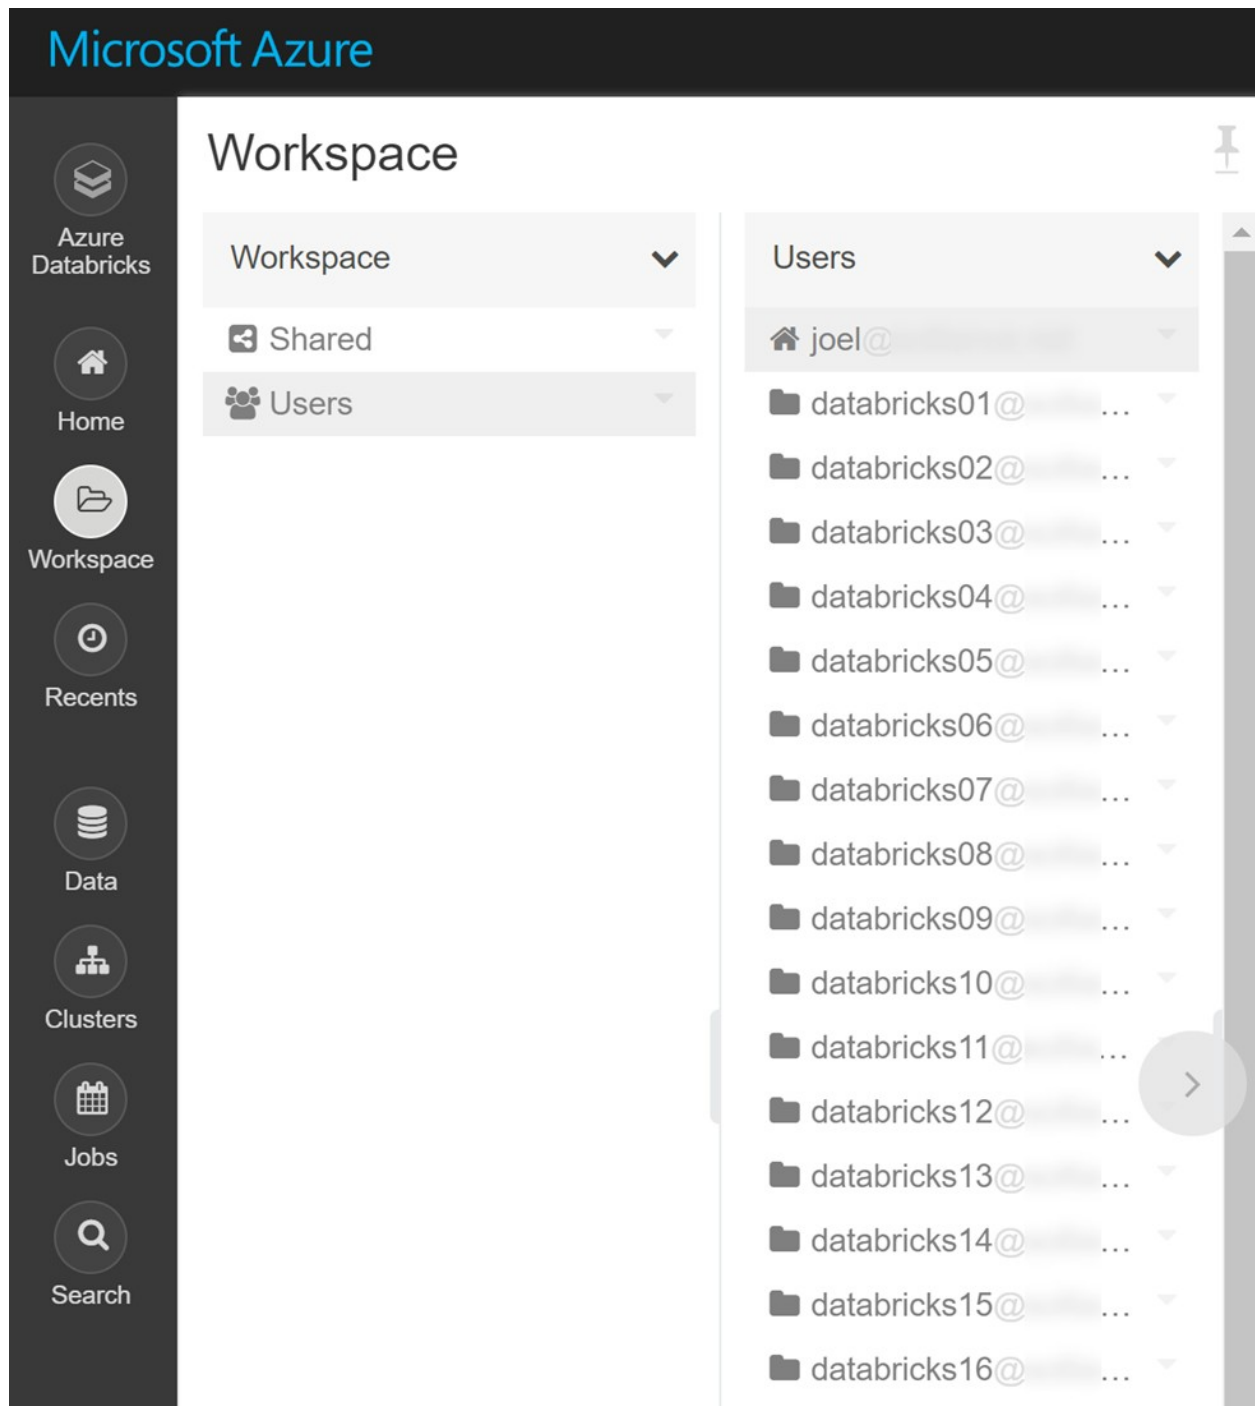
Create an Azure Databricks workspace and cluster

In this reading you can see the steps involved in the process of creating an Azure Databricks workspace and cluster.

When talking about the Azure Databricks workspace, we refer to two different things.

The first reference is the logical Azure Databricks environment in which clusters are created, data is stored (via DBFS), and in which the server resources are housed.

The second reference is the more common one used within the context of Azure Databricks. That is the special root folder for all of your organization's Databricks assets, including notebooks, libraries, and dashboards, as shown below:



The Databricks workspace folder is shown.

The first step to using Azure Databricks is to create and deploy a Databricks workspace, which is the logical environment. You can do this in the Azure portal.

Deploy an Azure Databricks workspace

1. Open the Azure portal

2. Click **Create a Resource** in the top left
3. Search for "Databricks"
4. Select *Azure Databricks*
5. On the Azure Databricks page select *Create*
6. Provide the required values to create your Azure Databricks workspace:
 - **Subscription**: Choose the Azure subscription in which to deploy the workspace.
 - **Resource Group**: Use **Create new** and provide a name for the new resource group.
 - **Location**: Select a location near you for deployment. For the list of regions that are supported by Azure Databricks, see [Azure services available by region](#).
 - **Workspace Name**: Provide a unique name for your workspace.
 - **Pricing Tier: Trial (Premium - 14 days Free DBUs)**. You must select this option when creating your workspace or you will be charged. The workspace will suspend automatically after 14 days. When the trial is over you can convert the workspace to **Premium** but then you will be charged for your usage.
7. Select **Review + Create**.
8. Select **Create**.

The workspace creation takes a few minutes. During workspace creation, the **Submitting deployment for Azure Databricks** tile appears on the right side of the portal. You might need to scroll right on your dashboard to view the tile. There's also a progress bar displayed near the top of the screen. You can watch either area for progress.

What is a cluster?

The notebooks are backed by clusters, or networked computers, that work together to process your data. The first step is to create a cluster.

Create a cluster

1. When your Azure Databricks workspace creation is complete, select the link to go to the resource.
2. Select **Launch Workspace** to open your Databricks workspace in a new tab.
3. In the left-hand menu of your Databricks workspace, select **Clusters**.
4. Select **Create Cluster** to add a new cluster.

Create Cluster

New Cluster

CancelCreate Cluster

0 Workers: 0.0 GB Memory, 0 Cores, 0 DBU

1 Driver: 14.0 GB Memory, 4 Cores, 0.75 DBU ?

Cluster Name



Cluster Mode ?



Pool ?



Databricks Runtime Version ?

[Learn more](#)New

This Runtime version supports only Python 3.

Autopilot Options

☒ Terminate after minutes of inactivity ?

Node Type ?

14.0 GB Memory, 4 Cores, 0.75 DBU



▶ Advanced Options

The create cluster page.

5. Enter a name for your cluster. Use your name or initials to easily differentiate your cluster from those of your co-workers.

6. Select the **Cluster Mode: Single Node**.

7. Select the **Databricks RuntimeVersion: Runtime: 7.3 LTS (Scala 2.12, Spark 3.0.1)**.

8. Under **Autopilot Options**, leave the box **checked** and in the text box enter **45**.

9. Select the **Node Type: Standard_DS3_v2**.

10. Select **Create Cluster**.